

R-MINI PROJECT

PROJECT TOPIC: Prediction and Analysis on Diabetes Dataset

Project Description:

The system will allow us to predict if the patient has diabetes on the basis of certain diagnostic measures available in the dataset. The different steps involved in EDA include: 1.Data Collection, 2.Data Cleaning and 3.Data Visualization. This project first conducts Exploratory Data Analysis (EDA) and data visualization on the diabetes dataset and then predicts the diabetes.

Exploratory Data Analysis (EDA)

1. Descriptive statistics

Attribute type, Class distribution, Mean, Standard Deviation, Median, Quartile, Skewness, and Correlation.

2. Data Visualization

Histogram plot

Density plot

Box and Whisker plot

Bar plot

Missing data map

Pairwise correlation plot

Prediction on Diabetes

We compare the performance for the following classifiers:

1. Logistic Regression

Logistic regression is basically a supervised classification algorithm. In a classification problem, the target variable (or output), y , can take only discrete values for given set of features (or inputs), X .

2. Support Vector Machine (SVM)

“Support Vector Machine” (SVM) is a supervised machine learning algorithm which can be used for both classification and regression challenges. However, it is mostly used in classification problems. In the SVM algorithm, we plot each data item as a point in n-dimensional space (where n is number of features you have) with the value of each feature being the value of a particular coordinate. Then, we perform classification by finding the hyper-plane that differentiates the two classes very well

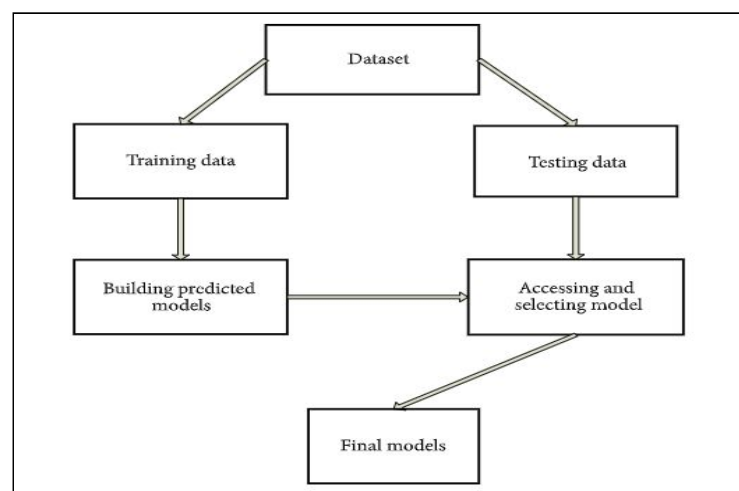
3. Random Forest

The random forest algorithm works by aggregating the predictions made by multiple decision trees of varying depth. Every decision tree in the forest is trained on a subset of the dataset called the bootstrapped dataset.

Dataset:

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	
1	id	chol	stab.glu	hdl	ratio	glyhb	location	age	gender	height	weight	frame	bp.1s	bp.1d	bp.2s	bp.2d	waist	hip	time.ppn	
2	1000	203	82	56	3.6	4.31	Buckingham	46	female	62	121	medium	118	59				29	38	720
3	1001	165	97	24	6.9	4.44	Buckingham	29	female	64	218	large	112	68				46	48	360
4	1002	228	92	37	6.2	4.64	Buckingham	58	female	61	256	large	190	92	185	92		49	57	180
5	1003	78	93	12	6.5	4.63	Buckingham	67	male	67	119	large	110	50				33	38	480
6	1005	249	90	28	8.9	7.72	Buckingham	64	male	68	183	medium	138	80				44	41	300
7	1008	248	94	69	3.6	4.81	Buckingham	34	male	71	190	large	132	86				36	42	195
8	1011	195	92	41	4.8	4.84	Buckingham	30	male	69	191	medium	161	112	161	112		46	49	720
9	1015	227	75	44	5.2	3.94	Buckingham	37	male	59	170	medium						34	39	1020
10	1016	177	87	49	3.6	4.84	Buckingham	45	male	69	166	large	160	80	128	86		34	40	300
11	1022	263	89	40	6.6	5.78	Buckingham	55	female	63	202	small	108	72				45	50	240
12	1024	242	82	54	4.5	4.77	Louisa	60	female	65	156	medium	130	90	130	90		39	45	300
13	1029	215	128	34	6.3	4.97	Louisa	38	female	58	195	medium	102	68				42	50	90
14	1030	238	75	36	6.6	4.47	Louisa	27	female	60	170	medium	130	80				35	41	720
15	1031	183	79	46	4	4.59	Louisa	40	female	59	165	medium						37	43	60
16	1035	191	76	30	6.4	4.67	Louisa	36	male	69	183	medium	100	66				36	40	225
17	1036	213	83	47	4.5	3.41	Louisa	33	female	65	157	medium	130	90	120	96		37	41	240
18	1037	255	78	38	6.7	4.33	Louisa	50	female	65	183	medium	130	100				37	43	180
19	1041	230	112	64	3.6	4.53	Louisa	20	male	67	159	medium	100	90				31	39	1440
20	1045	194	81	36	5.4	5.28	Louisa	36	male	64	126	medium	110	76				30	34	120
21	1250	196	206	41	4.8	11.24	Buckingham	62	female	65	196	large	178	90				46	51	540
22	1252	186	97	50	3.7	6.49	Buckingham	70	male	67	178	large	148	88	148	84		42	41	1020

Our research dataset is divided into two parts; two-thirds of the data is used as a training set, and one-third of the dataset is defined as a testing set to evaluate the performance of several classifiers. All classifiers were fitted to the same training and testing data. The specific process is: .



Libraries:

- **Randomforest**

Random Forest implements Breiman's random forest algorithm (based on Breiman's and Cutler's original FORTRAN code) for classification and regression. It can also be used in unsupervised mode for assessing proximities among data points.

- **Caret**

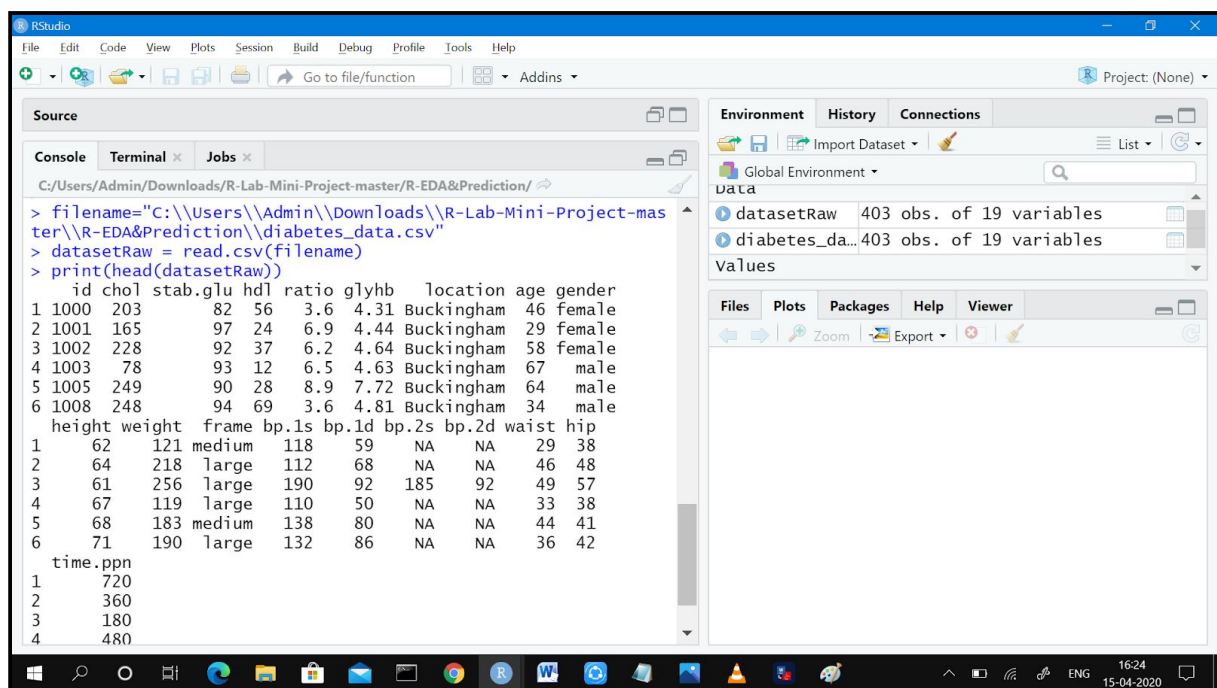
The caret package (short for Classification And REgression Training) is a set of functions that attempt to streamline the process for creating predictive models. The package contains tools for:

- data splitting
- pre-processing
- feature selection
- model tuning using resampling
- variable importance estimation
- as well as other functionality.

Output:

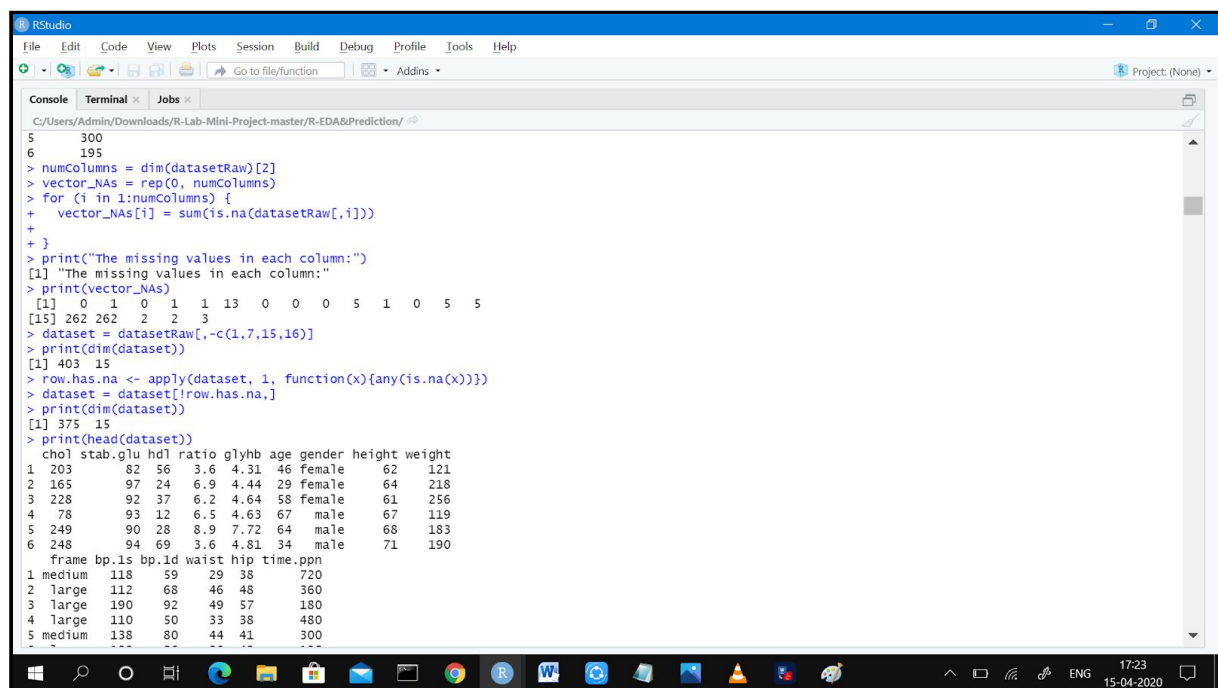
Step1: Load the data

```
> filename="C:\\Users\\Admin\\Desktop\\R-Lab-Mini-Project \\diabetes_data.csv"
> datasetRaw = read.csv(filename)
> print(head(datasetRaw))
```



Step 2: Clean the data

```
> numColumns = dim(datasetRaw)[2]
> vector_NAs = rep(0, numColumns)
> for (i in 1:numColumns) {
+   vector_NAs[i] = sum(is.na(datasetRaw[,i]))
+ }
> print("The missing values in each column:")
> print(vector_NAs)
# deletes columns 15 and 16 due to many missing values
# deletes column 1 (id), column 7 (location) because they contain no useful information
> print(dim(dataset))
remove the row with missing values
> row.has.na <- apply(dataset, 1, function(x){any(is.na(x))})
> dataset = dataset[!row.has.na,]
> print(dim(dataset))
> print(head(dataset))
```



The screenshot shows the RStudio interface with the console window open. The console displays the execution of the R code from the previous block. The output shows the dimensions of the dataset after removing columns 1, 7, 15, and 16, and rows with missing values. The head of the dataset is printed, showing the first 5 rows and columns 2 through 14. The data includes variables like cholesterol, blood sugar, hemoglobin, and body mass index.

```
5      300
6      195
> numColumns = dim(datasetRaw)[2]
> vector_NAs = rep(0, numColumns)
> for (i in 1:numColumns) {
+   vector_NAs[i] = sum(is.na(datasetRaw[,i]))
+ }
> print("The missing values in each column:")
[1] "The missing values in each column:"
> print(vector_NAs)
[1] 0 1 0 1 1 13 0 0 0 5 1 0 5 5
[15] 262 262 2 2 3
> dataset = datasetRaw[,-c(1,7,15,16)]
> print(dim(dataset))
[1] 403 15
> row.has.na <- apply(dataset, 1, function(x){any(is.na(x))})
> dataset = dataset[!row.has.na,]
> print(dim(dataset))
[1] 375 15
> print(head(dataset))
      chol stab.glu hdl ratio glyhb age gender height weight
1    203    82  56   3.6  4.31  46 female    62    121
2    165    97  24   6.9  4.44  29 female    64    218
3    228    92  37   6.2  4.64  58 female    61    256
4     78    93  12   6.5  4.63  67 male     67    119
5    249    90  28   8.9  7.72  64 male     68    183
6    248    94  69   3.6  4.81  34 male     71    190
      frame bp.ls bp.id waist hip time.ppn
1 medium   118    59    29  38    720
2 large   112    68    46  48    360
3 large   190    92    49  57    180
4 large   110    50    33  38    480
5 medium   138    80    44  41    300
```

encodes the class label (column 5): Glycosylated haemoglobin > 7.0 is taken as a positive diagnosis of diabetes.

```
> dataset[,5] = if else(dataset[,5] >= 7.0, 1, 0)
> dataset[,5] = factor(dataset[,5])
```

encode the categorical data (column-7 gender)

```
> dataset[,7] = if else(dataset[,7] == "female", 0, 1)
> dataset[,7] = factor(dataset[,7])
```

encode the categorical data (column-10 frame)

```

> dataset[,10] = if else(dataset[,10] == "small", 0, if else(dataset[,10] == "medium", 1,2) )
> dataset[,10] = factor(dataset[,10])
# Descriptive statistics %%%%%%%%%
# display the first 20 rows
> print(head(dataset, n=20))

```

```

RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/
Console Terminal Jobs
chol stab glu hdl ratio glyhb age gender height weight
1 203 82 56 3.6 0 46 0 62 121
2 165 97 24 6.9 0 29 0 64 218
3 228 92 37 6.2 0 58 0 61 256
4 78 93 12 6.5 0 67 1 67 119
5 249 90 28 8.9 1 64 1 68 183
6 248 94 69 3.6 0 34 1 71 190
7 195 92 41 4.8 0 30 1 69 191
9 177 87 49 3.6 0 45 1 69 166
10 263 89 40 6.6 0 55 0 63 202
11 242 82 54 4.5 0 60 0 65 156
12 215 128 34 6.3 0 38 0 58 195
13 238 75 36 6.6 0 27 0 60 170
15 191 76 30 6.4 0 36 1 69 183
16 213 83 47 4.5 0 33 0 65 157
17 255 78 38 6.7 0 50 0 65 183
18 230 112 64 3.6 0 20 1 67 159
19 194 81 36 5.4 0 36 1 64 126
20 196 206 41 4.8 1 62 0 65 196
21 186 97 50 3.7 0 70 1 67 178
22 234 65 76 3.1 0 47 1 67 230
frame bp.ls bp.id waist hip time.ppn
1 1 118 59 29 38 720
2 2 112 68 46 48 360
3 2 190 92 49 57 180
4 2 110 50 33 38 480
5 1 138 80 44 41 300
6 2 132 86 36 42 195
7 1 161 112 46 49 720
9 2 160 80 34 40 300
10 0 108 72 45 50 240
11 1 130 90 39 45 300
12 1 102 68 42 50 90
13 1 130 80 35 41 720

```

```

# display the dimensions of the dataset
> print(dim(dataset))
[1] 375 15
> print(head(dataset))

```

```

RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/
Console Terminal Jobs
21 2 148 88 42 41 1020
22 2 137 100 45 46 480
> print(dim(dataset))
[1] 375 15
> print(sapply(dataset,class))
chol stab glu hdl ratio glyhb age
"integer" "integer" "integer" "numeric" "factor" "integer"
gender height weight frame bp.ls bp.id
"factor" "integer" "integer" "factor" "integer" "integer"
waist hip time.ppn
"integer" "integer" "integer"
> y = dataset$glyhb
> print(cbind(freq=table(y),percentage=prop.table(table(y))*100))
freq percentage
0 317 84.53333
1 58 15.46667
> print(table(y))
y
0 1
317 58
> print(table(y)/length(y))
y
0 1
0.8453333 0.1546667
> print(summary(dataset))
chol stab glu hdl
Min. : 78.0 Min. : 48.0 Min. : 12.00
1st Qu.:179.0 1st Qu.: 81.0 1st Qu.: 38.00
Median :204.0 Median : 90.0 Median : 46.00
Mean :207.6 Mean :107.6 Mean : 50.43
3rd Qu.:229.5 3rd Qu.:108.5 3rd Qu.: 59.00
Max. :443.0 Max. :385.0 Max. :120.00
ratio glyhb age gender
Min. : 1.500 0:317 Min. :19.00 0:220

```

```
> # encode the class label (column 5): Glycosylated hemoglobin > 7.0 is taken as a positive diagnosis of diabetes.
```

```
> dataset[,5] = ifelse(dataset[,5] >= 7.0, 1, 0)
```

```
> dataset[,5] = factor(dataset[,5])
```

```
> # encode the categorical data (column-7 gender)
```

```
> dataset[,7] = ifelse(dataset[,7] == "female", 0, 1)
```

```
> dataset[,7] = factor(dataset[,7])
```

```
> # encode the categorical data (column-10 frame)
```

```
> dataset[,10] = ifelse(dataset[,10] == "small", 0, ifelse(dataset[,10] == "medium", 1, 2) )
```

```
> dataset[,10] = factor(dataset[,10])
```

```
> # Descriptive statistics %%%%%%%%%
```

```
> # display the first 20 rows
```

```
> print(head(dataset, n=20))
```

The screenshot shows the RStudio interface with the following code in the console:

```
> dataset[,7] = ifelse(dataset[,7] == "female", 0, 1)
> dataset[,7] = factor(dataset[,7])
> dataset[,10] = ifelse(dataset[,10] == "small", 0, ifelse(dataset[,10] == "medium", 1, 2) )
> dataset[,10] = factor(dataset[,10])
> print(head(dataset, n=20))
```

The output shows the first 20 rows of the dataset with columns: chol, stab.glu, hdl, ratio, glyhb, age, gender, height, weight, frame, bp.1s, waist, hip, time.ppn. The data is as follows:

	chol	stab.glu	hdl	ratio	glyhb	age	gender	height	weight	frame	bp.1s	waist	hip	time.ppn
1	203	82	56	3.6	0	46	0	62	121					
2	165	97	34	6.9	0	29	0	64	218					
3	228	92	37	6.2	0	58	0	61	256					
4	78	93	12	6.5	0	67	1	67	119					
5	249	90	28	8.9	1	64	1	68	183					
6	248	94	69	3.6	0	34	1	71	190					
7	195	92	41	4.8	0	30	1	69	191					
9	177	87	49	3.6	0	45	1	69	166					
10	263	89	40	6.6	0	55	0	63	202					
11	242	82	54	4.5	0	60	0	65	156					
12	215	128	34	6.3	0	38	0	58	195					
13	258	75	56	6.6	0	27	0	60	170					
15	191	76	30	6.4	0	26	1	69	183					
16	213	83	47	4.5	0	53	0	65	157					
17	255	78	38	6.7	0	50	0	65	183					
18	230	112	64	3.6	0	20	1	67	159					
19	194	81	36	5.4	0	36	1	64	126					
20	196	206	41	4.8	1	62	0	65	196					
22	146	97	50	3.7	0	70	1	67	178					
22	234	65	76	3.1	0	47	1	67	230					

The dimensions of the dataset are shown as:

```
[1] 375 15
```

```
> # display the dimensions of the dataset
```

```
> print(dim(dataset))
```

```
[1] 375 15
```

```
> # list types for each attribute
```

```
> print(sapply(dataset, class))
```

```
chol stab.glu hdl ratio glyhb age gender height weight frame bp.1s
"integer" "integer" "integer" "numeric" "factor" "integer" "factor" "integer" "integer" "factor"
"integer"
bp.1d waist hip time.ppn
"integer" "integer" "integer" "integer"
```

```
> # distribution of the class labels
> y = dataset$glyhb
> print(cbind(freq=table(y),percentage=prop.table(table(y))*100))
```

```
freq percentage
0 317 84.53333
1 58 15.46667
```

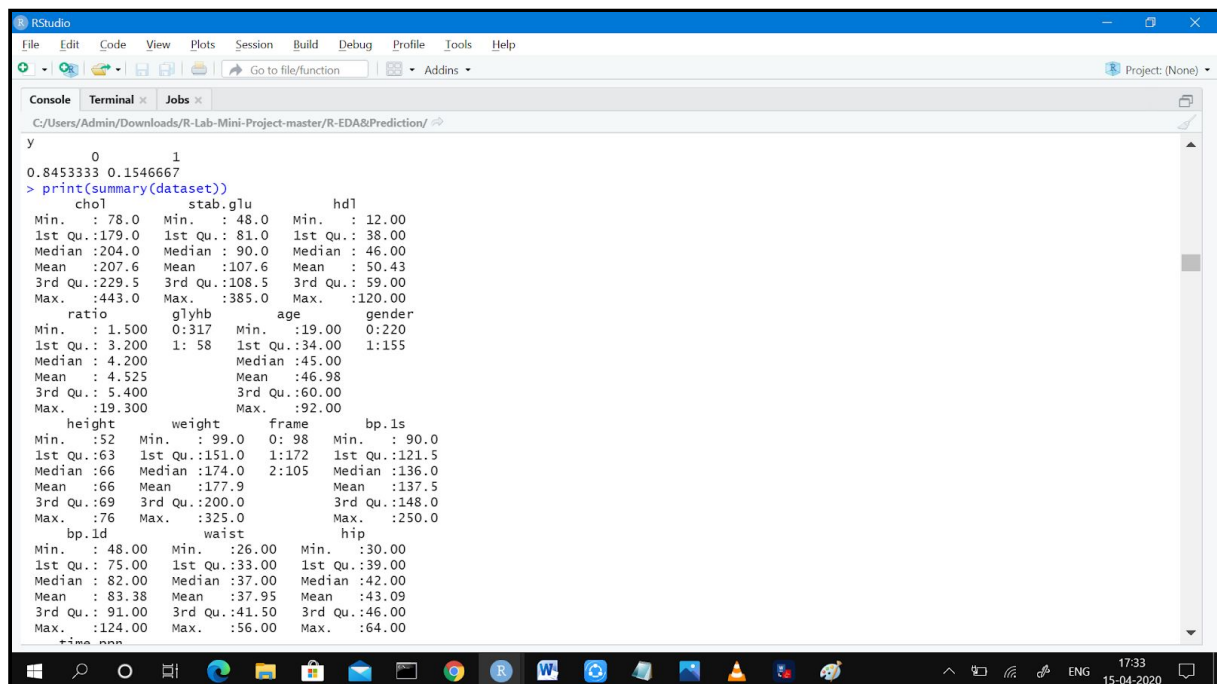
```
> print(table(y))
```

```
y
0 1
317 58
```

```
> print(table(y)/length(y))
```

```
y
0 1
0.8453333 0.1546667
```

```
> # summarize the dataset
> print(summary(dataset))
```



```
RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins
Project: (None)
Console Terminal Jobs
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/
y
0 1
0.8453333 0.1546667
> print(summary(dataset))
 chol      stab.glu    hdl      ratio      age      height      weight      bp.ls
Min.   : 78.0   Min.   : 48.0   Min.   : 12.00
1st Qu.:179.0   1st Qu.: 81.0   1st Qu.: 38.00
Median :204.0   Median : 90.0   Median : 46.00
Mean   :207.6   Mean   :107.6   Mean   : 50.43
3rd Qu.:229.5   3rd Qu.:108.5   3rd Qu.: 59.00
Max.   :443.0   Max.   :385.0   Max.   :120.00
 ratio      glyhb      age      gender
Min.   : 1.500   0:317   Min.   :19.00   0:220
1st Qu.: 3.200   1: 58   1st Qu.:34.00   1:155
Median : 4.200   Median :45.00
Mean   : 4.525   Mean   :46.98
3rd Qu.: 5.400   3rd Qu.:60.00
Max.   :19.300   Max.   :92.00
 height      weight      frame      bp.ls
Min.   :52   Min.   :99.0   0: 98   Min.   : 90.0
1st Qu.:63   1st Qu.:151.0   1:172   1st Qu.:121.5
Median :66   Median :174.0   2:105   Median :136.0
Mean   :66   Mean   :177.9   Mean   :137.5
3rd Qu.:69   3rd Qu.:200.0   3rd Qu.:148.0
Max.   :76   Max.   :325.0   Max.   :250.0
 bp.la      waist      hip
Min.   :48.00   Min.   :26.00   Min.   :30.00
1st Qu.:75.00   1st Qu.:33.00   1st Qu.:39.00
Median :82.00   Median :37.00   Median :42.00
Mean   :83.38   Mean   :37.95   Mean   :43.09
3rd Qu.:91.00   3rd Qu.:41.50   3rd Qu.:46.00
Max.   :124.00   Max.   :56.00   Max.   :64.00
 time.ppn
```

Step 3: To Do EDA (Analysis of the dataset)

```
# Standard Deviations for the non-categorical columns
```

```
> std=sapply(dataset[, -c(5,7,10)],sd)
> print('The standard deviations are:')
[1] "The standard deviations are:"
```

```
> print(std)
```

```
chol stab.glu hdl ratio age height weight bp.ls
44.700780 54.082496 17.444346 1.755499 16.661203 3.915210 40.568940 23.178154
 bp.la waist hip time.ppn
13.544167 5.777105 5.642679 309.056806
```


Skewness

#The further the distribution of the skew value from zero,

the larger the skew to the left (negative skew value) or right (positive skew value).

```
> library(e1071) # the library for skewness
```

```
> skew=apply(dataset[,c(5,7,10)], 2, skewness)
```

```
> print(skew)
```

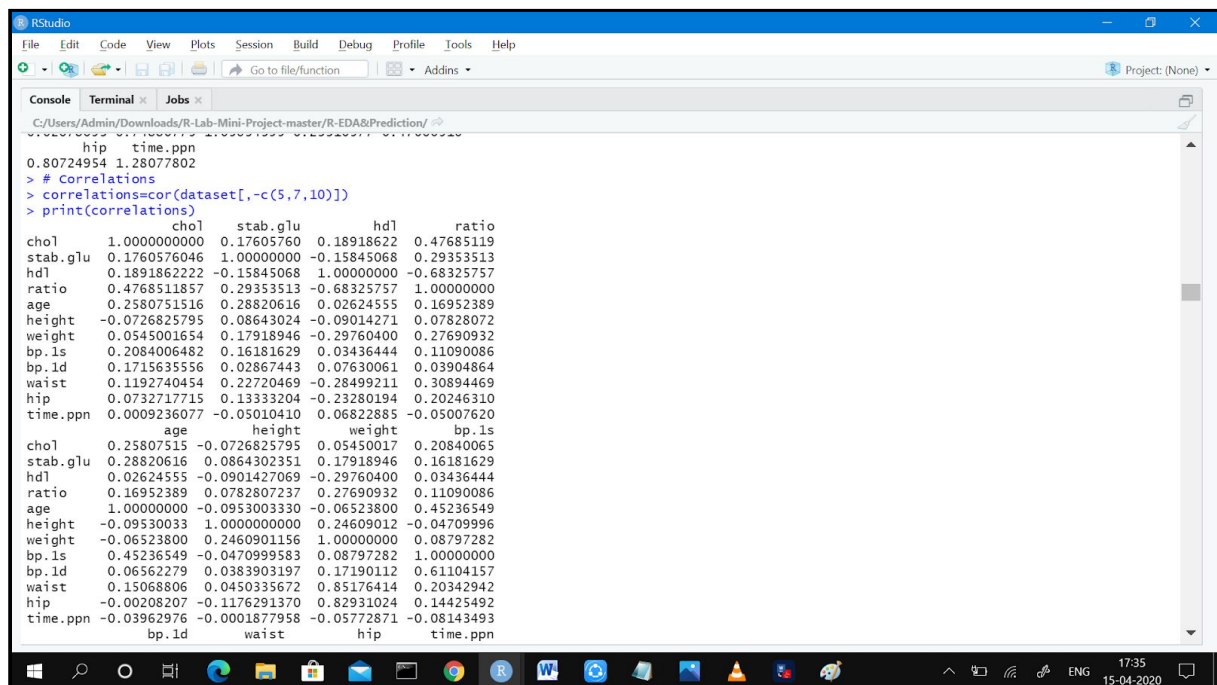
```
chol stab.glu hdl ratio age height weight bp.1s
0.97739823 2.69790949 1.21275829 2.24132546 0.30061280 0.02678693 0.74880775 1.05634395
bp.1d waist hip time.ppn
0.23310577 0.47060516 0.80724954 1.28077802
```

CORRELATION:

Correlations

```
> correlations=cor(dataset[,c(5,7,10)])
```

```
> print(correlations)
```

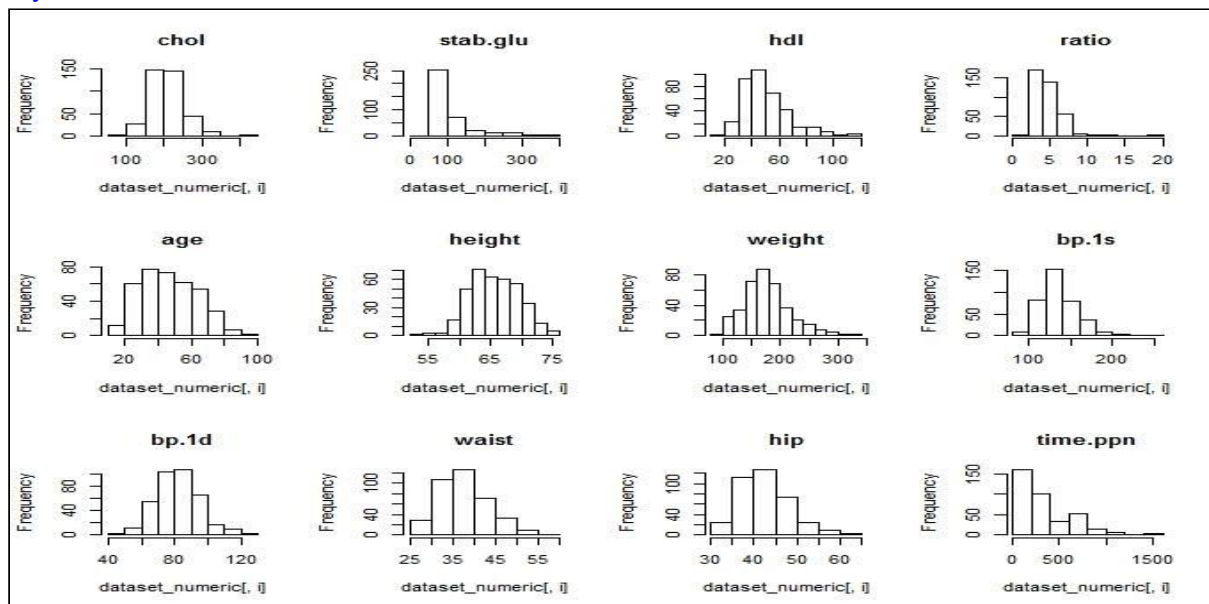


```
RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins Project: (None)
Console Terminal Jobs
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/
hip time.ppn
0.80724954 1.28077802
> # Correlations
> correlations=cor(dataset[,c(5,7,10)])
> print(correlations)
      chol      stab.glu      hdl      ratio
chol  1.0000000000  0.17605760  0.18918622  0.47685119
stab.glu 0.1760576046  1.00000000 -0.15845068  0.29353513
hdl      0.1891862222 -0.15845068  1.00000000 -0.68325757
ratio    0.4768511857  0.29353513 -0.68325757  1.00000000
age      0.2580751516  0.28820616  0.02624555  0.16952389
height   -0.0726825795  0.08643024 -0.09014271  0.07828072
weight   0.0545001654  0.17918946 -0.29760400  0.27690932
bp.1s    0.2084006482  0.16181629  0.03436444  0.11090086
bp.1d    0.1715635356  0.02867443  0.07630061  0.03904864
waist    0.1192740454  0.22720469 -0.28499211  0.30894469
hip      0.0732717715  0.13333204 -0.23280194  0.20246310
time.ppn 0.0009236077 -0.05010410  0.06822885 -0.05007620
      age      height      weight      bp.1s
chol  0.25807515 -0.0726825795  0.05450017  0.20840065
stab.glu 0.28820616  0.0864302351  0.17918946  0.16181629
hdl      0.02624555 -0.0901427069 -0.29760400  0.03436444
ratio    0.16952389  0.0782807237  0.27690932  0.11090086
age      1.00000000 -0.0953003330 -0.06523800  0.45236549
height   -0.09530033  1.0000000000  0.24609012 -0.04709996
weight   -0.06523800  0.2460901156  1.00000000  0.08797282
bp.1s    0.45236549 -0.0470999583  0.08797282  1.00000000
bp.1d    0.06562279  0.0383903197  0.17190112  0.61104157
waist    0.15068806  0.0450335672  0.85176414  0.20342942
hip      -0.00208207 -0.1176291370  0.82931024  0.14425492
time.ppn -0.03962976 -0.0001877958 -0.05772871 -0.08143493
      bp.1d      waist      hip      time.ppn
```


PART 2 : Data visualizations:


1:Histogram

```
> dataset_numeric = dataset[,-c(5,7,10)]  
> #Histograms  
> par(mfrow=c(3,4)) # put four figures in a row (2*4)  
> for (i in 1:12) {  
+   hist(dataset_numeric[,i],main=names(dataset_numeric)[i])  
+ }
```



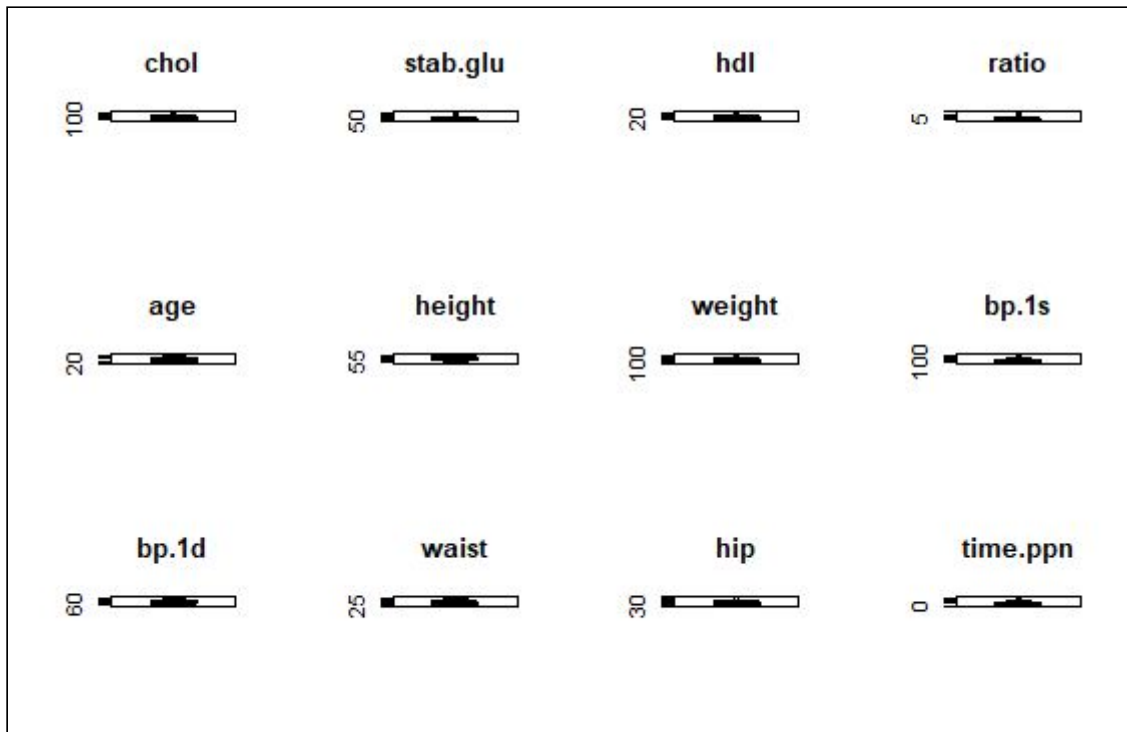
2:Density Plots

```
> par(mfrow=c(3,4))  
> for(i in 1:12) {  
+   plot(density(dataset_numeric[,i]), main=names(dataset_numeric)[i])  
+ }
```



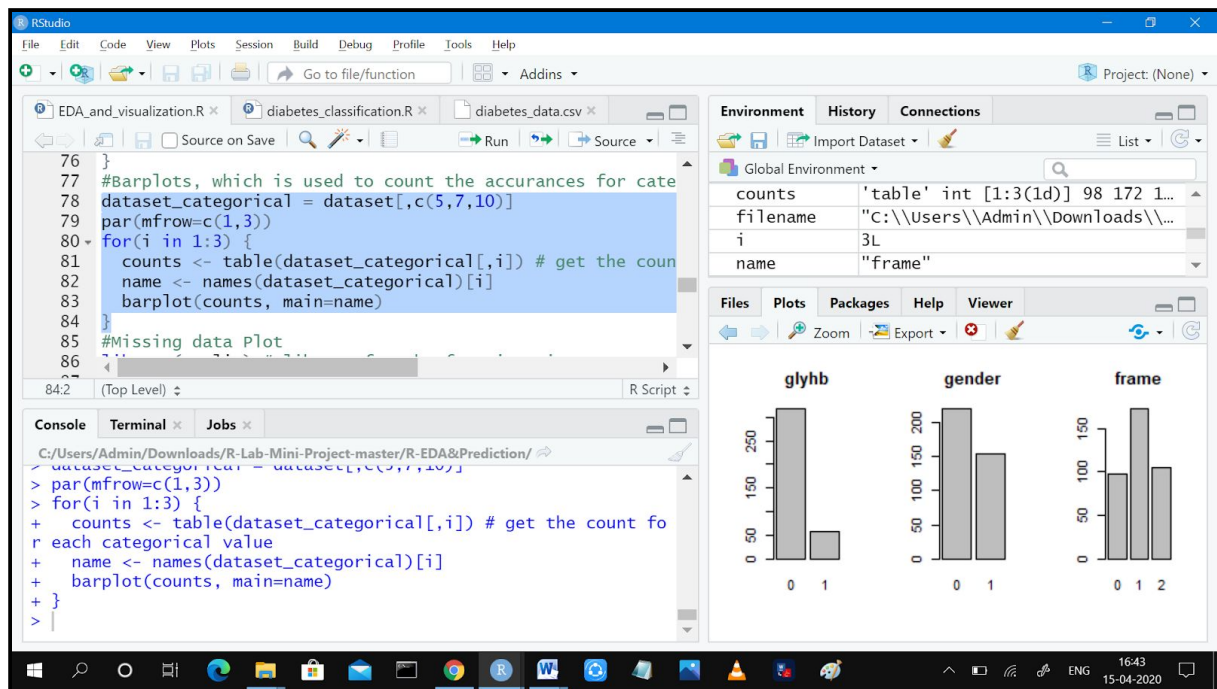
3:Box And Whisker Plots

```
> par(mfrow=c(3,4))  
> for(i in 1:12) {  
+   boxplot(dataset_numeric[,i], main=names(dataset_numeric)[i])  
+ }
```



4:Barplots:

```
> dataset_categorical = dataset[,c(5,7,10)]  
> par(mfrow=c(1,3))  
> for(i in 1:3) {  
+   counts <- table(dataset_categorical[,i]) # get the count for each categorical value  
+   name <- names(dataset_categorical)[i]  
+   barplot(counts, main=name)  
+ }
```

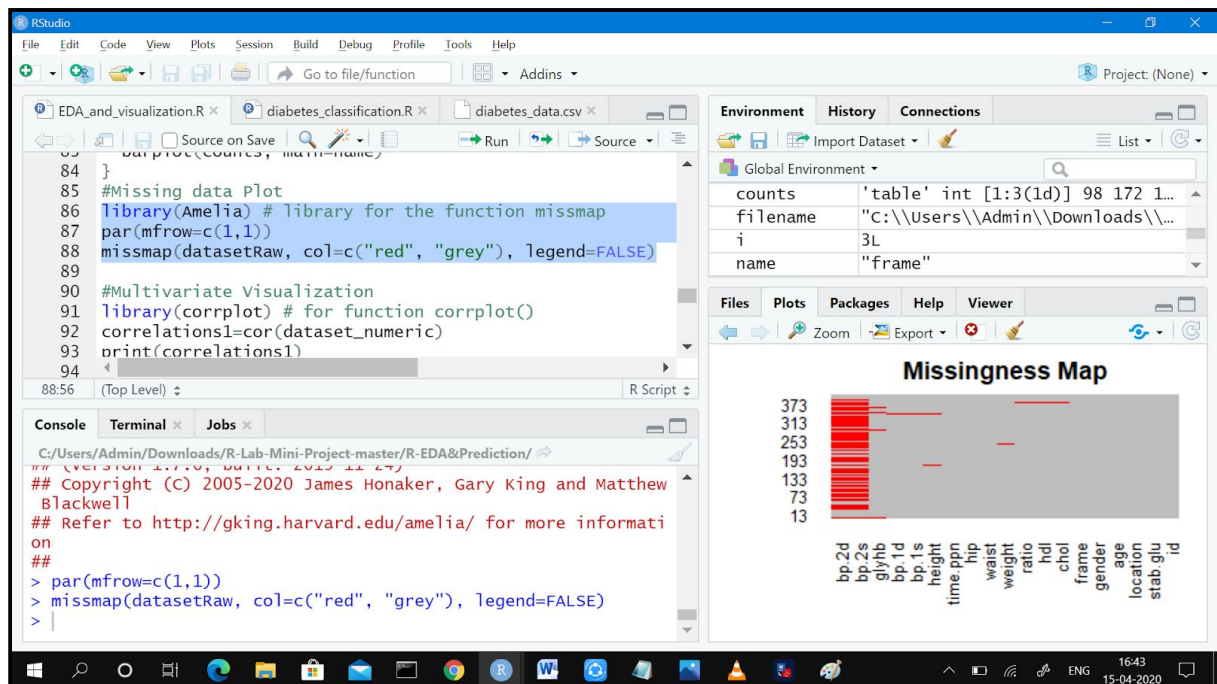


5:Missing data Plot

```

> library (Amelia) # library for the function missmap
> par(mfrow=c(1,1))
> missmap (datasetRaw, col=c("red", "grey"), legend=FALSE)

```



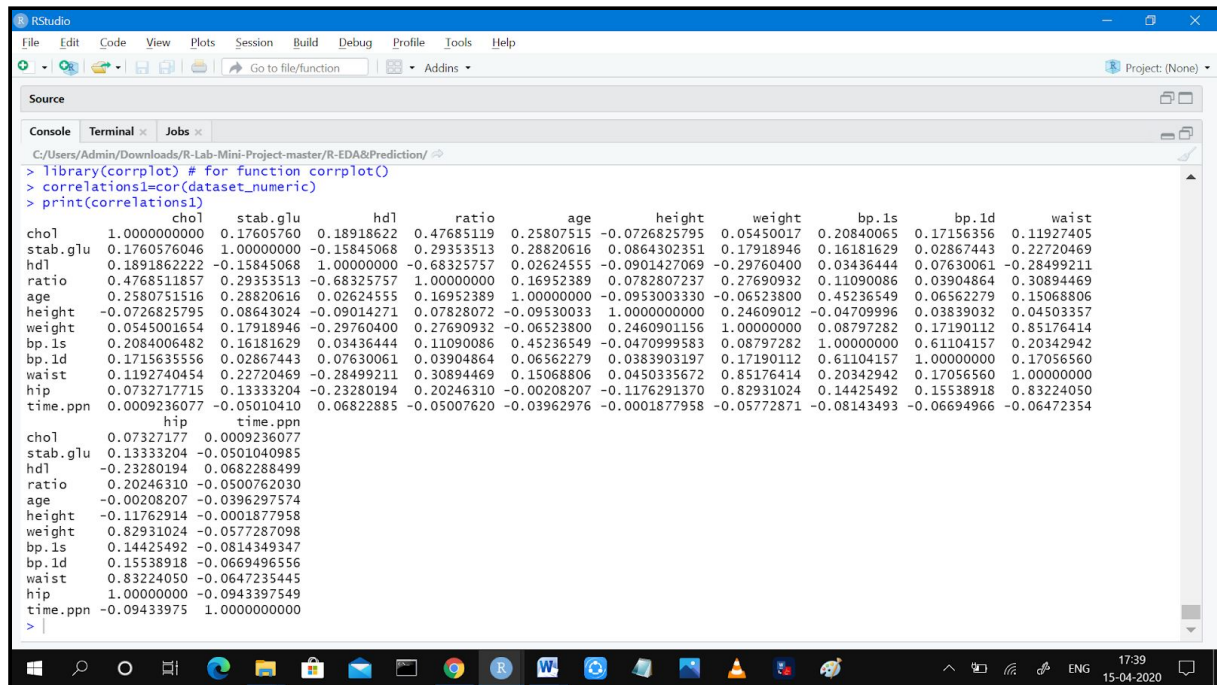
6: Multivariate Visualization

```
> library(corrplot) # for function corrplot()
```

Corrplot 0.84 loaded

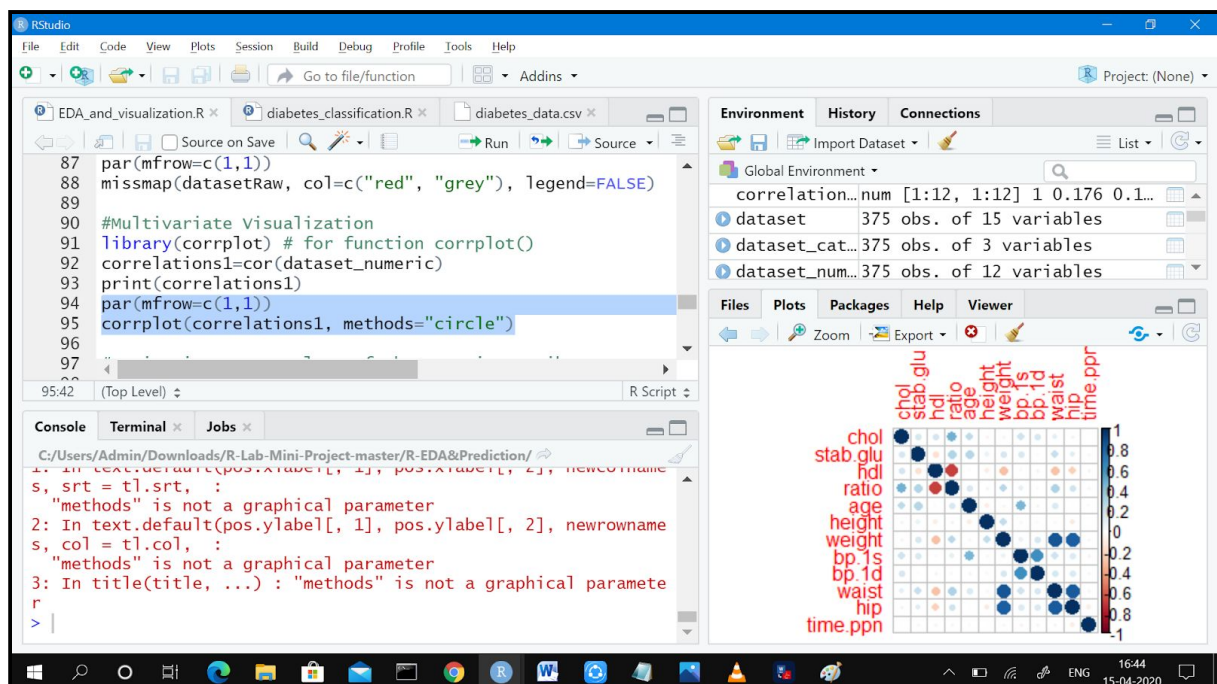
```
> correlations1=cor (dataset_numeric)
```

```
> print (correlations1)
```




```
> par(mfrow=c(1,1))
```

```
> corrplot(correlations1, methods="circle")
```



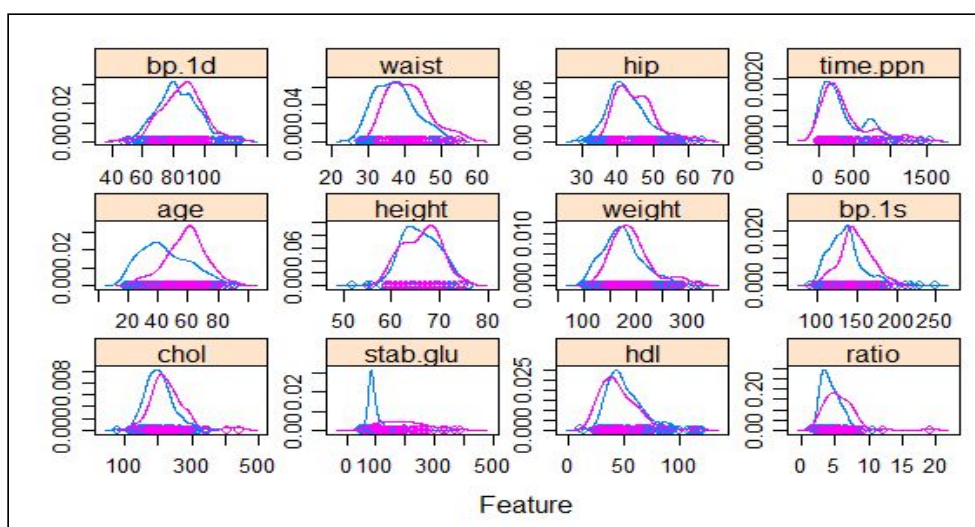
7: Pairwise scatterplots of the numeric attributes

```
> par(mfrow=c(1,1))
pairs (dataset_numeric)
> #Scatterplot Matrix By Class (use different color to distinguish different class)
> par(mfrow=c(1,1))
> pairs (dataset_numeric, col=dataset[,5])
```



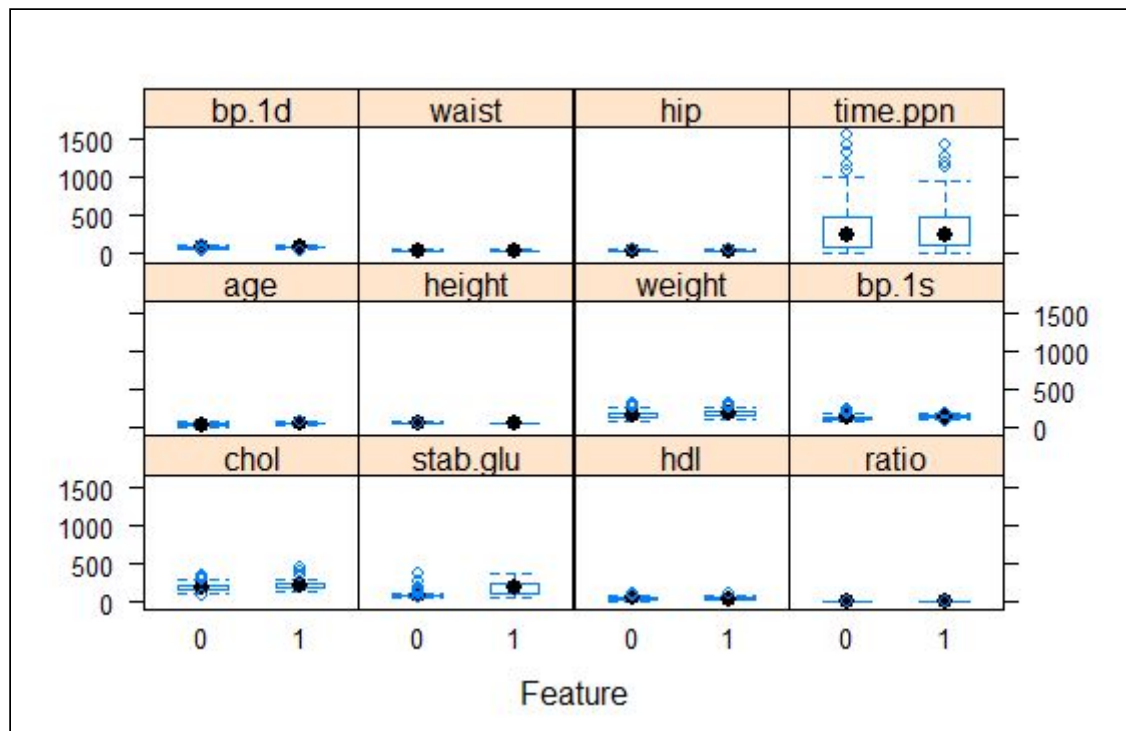
8: Density By Class

```
> library(caret)
> # load the data
> data(iris)
> # density plots for each attribute by class value
> x <- dataset_numeric
> y <- dataset[,5]
> scales <- list(x=list(relation="free"), y=list(relation="free"))
> par(mfrow=c(1,1))
> feature Plot(x=dataset_numeric, y=dataset[,5], plot="density", scales=scales)
```



9:Box And Whisker Plots By Class

```
> featurePlot(x=dataset_numeric, y=dataset[,5], plot="box")
```



```
# Load libraries
```

```
> library(randomForest)
```

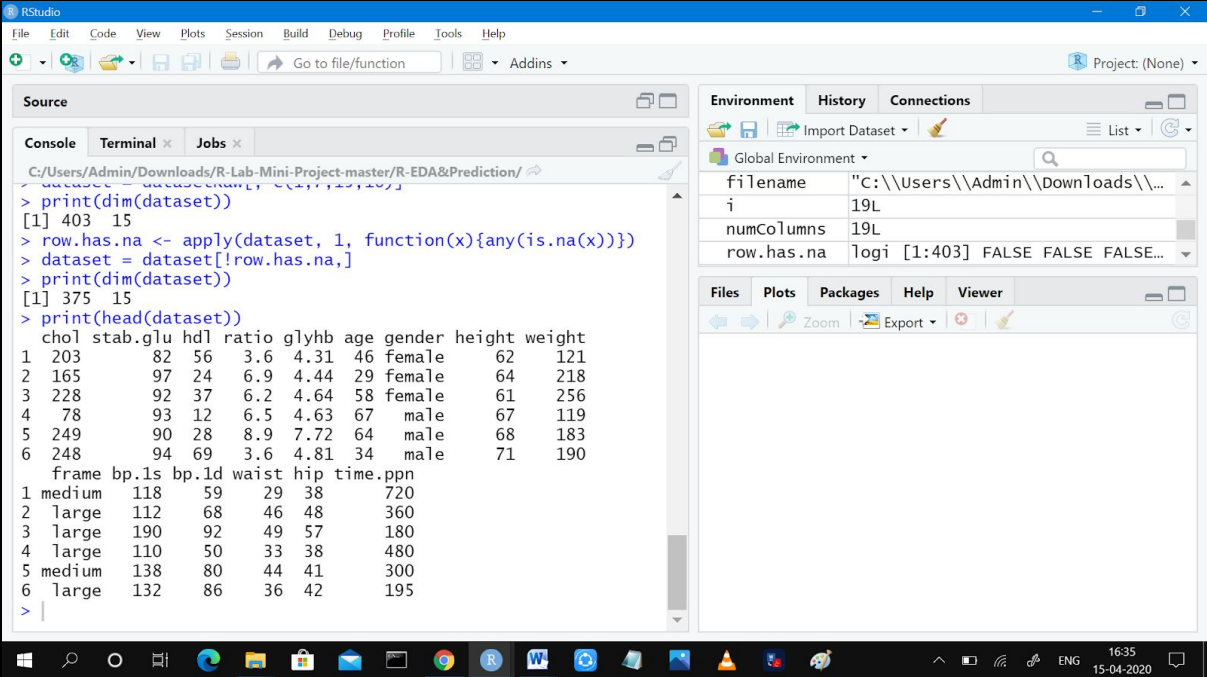
```
> library(caret)
```

```
> # load the data
```

```
> filename="C:\\Users\\Admin \\Desktop\\ R-Lab-Mini-Project \\diabetes_data.csv"
```

```
> datasetRaw = read.csv(filename)
```

```
> print(head(datasetRaw))
```



```
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/
> dataset = datasetRaw[, 1:(7+4+4)]
> print(dim(dataset))
[1] 403 15
> row.has.na <- apply(dataset, 1, function(x){any(is.na(x))})
> dataset = dataset[!row.has.na,]
> print(dim(dataset))
[1] 375 15
> print(head(dataset))
  chol stab.glu hdl ratio glyhb age gender height weight
1 203      82  56  3.6  4.31  46 female    62    121
2 165      97  24  6.9  4.44  29 female    64    218
3 228      92  37  6.2  4.64  58 female    61    256
4  78      93  12  6.5  4.63  67  male    67    119
5 249      90  28  8.9  7.72  64  male    68    183
6 248      94  69  3.6  4.81  34  male    71    190
  frame bp.ls bp.ld waist hip time.ppn
1 medium  118   59   29  38    720
2 large  112   68   46  48    360
3 large  190   92   49  57    180
4 large  110   50   33  38    480
5 medium 138   80   44  41    300
6 large  132   86   36  42    195
>
```

```
# clean the data
```

```
> numColumns = dim(datasetRaw)[2]
```

```
> vector_NAs = rep (0, numColumns)
```

```
> for (i in 1:numColumns) {
```

```
+ vector_NAs[i] = sum (is.na (datasetRaw [,i]))
```

```
+ 
```

```
+ }
```

```
> print ("The missing values in each column :")
```

```
[1] "The missing values in each column:"
```

```
> print (vector_NAs)
```

```
[1] 0 1 0 1 1 13 0 0 0 5 1 0 5 5 262 262 2 2 3
```

```
# delete columns 15 and 16 due to many missing values
```

```
> # delete column 1 (id), column 7 (location) because they contain no useful information
```

```
> dataset = datasetRaw[,-c(1,7,15,16)]
```

```
> print(dim(dataset))
```

```
[1] 403 15
```

```
# remove the row with missing values
```

```
> row.has.na <- apply(dataset, 1, function(x){any(is.na(x))})
```



```
> dataset = dataset[!row.has.na,]
> print(dim(dataset))
```

```
[1] 375 15
```

```
> print(head(dataset))
```

	ld	chol	stab	glu	hdl	ratio	glyhb	age	gender	height	weight	frame	bp.1s	bp.1d	waist	hip
1	203	82	56	3.6	4.31	46	female	62	121	medium	118	59	29	38		
2	165	97	24	6.9	4.44	29	female	64	218	large	112	68	46	48		
3	228	92	37	6.2	4.64	58	female	61	256	large	190	92	49	57		
4	78	93	12	6.5	4.63	67	male	67	119	large	110	50	33	38		
5	249	90	28	8.9	7.72	64	male	68	183	medium	138	80	44	41		
6	248	94	69	3.6	4.81	34	male	71	190	large	132	86	36	42		

```
# encode the class label (column 5): Glycosylated hemoglobin > 7.0 is taken as  
a positive diagnosis of diabetes.
```

```
> dataset[,5] = ifelse(dataset[,5] >= 7.0, 1, 0)
> dataset[,5] = factor(dataset[,5])
```

```
> # encode the categorical data (column-7 gender)
> dataset[,7] = ifelse(dataset[,7] == "female", 0, 1)
> dataset[,7] = factor(dataset[,7])
```

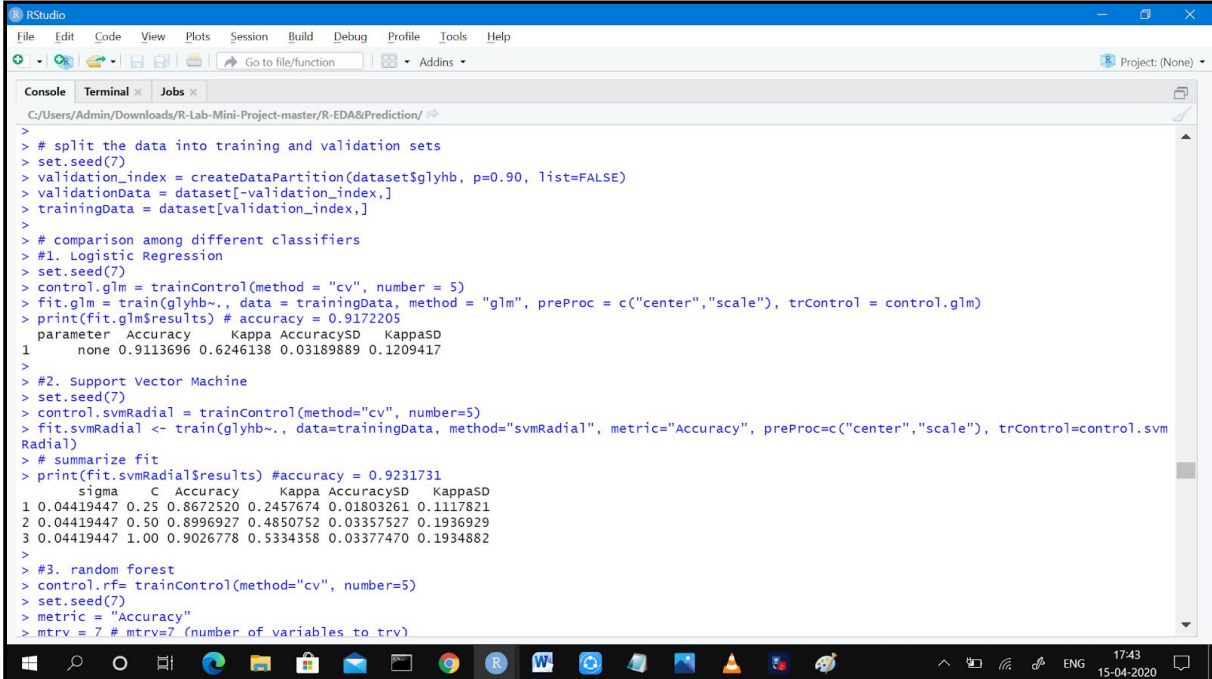
```
> # encode the categorical data (column-10 frame)
> dataset[,10] = ifelse(dataset[,10] == "small", 0, ifelse(dataset[,10] ==  
"medium", 1,2) )
> dataset[,10] = factor(dataset[,10])
```

```
> # split the data into training and validation sets
> set.seed(7)
> validation_index = createDataPartition(dataset$glyhb, p=0.90, list=FALSE)
> validationData = dataset[-validation_index,]
> trainingData = dataset[validation_index,]
```

PART 3: DATA CLASSIFICATION

#1. Logistic Regression

```
> set.seed(7)
> control.glm = trainControl(method = "cv", number = 5)
> fit.glm = train(glyhb~, data = trainingData, method = "glm", preProc = c("center","scale"),
trControl = control.glm)
> print(fit.glm$results)
```

The screenshot shows the RStudio interface with the console window open. The console displays the following R code and its output:

```
> # split the data into training and validation sets
> set.seed(7)
> validation_index = createDataPartition(dataset$glyhb, p=0.90, list=FALSE)
> validationData = dataset[-validation_index,]
> trainingData = dataset[validation_index,]
>
> # comparison among different classifiers
> #1. Logistic Regression
> set.seed(7)
> control.glm = trainControl(method = "cv", number = 5)
> fit.glm = train(glyhb~, data = trainingData, method = "glm", preProc = c("center","scale"), trControl = control.glm)
> print(fit.glm$results) # accuracy = 0.9172205
  parameter Accuracy Kappa AccuracySD KappaSD
1      none 0.9113696 0.6246138 0.03189889 0.1209417
>
> #2. Support Vector Machine
> set.seed(7)
> control.svmRadial = trainControl(method="cv", number=5)
> fit.svmRadial <- train(glyhb~, data=trainingData, method="svmRadial", metric="Accuracy", preProc=c("center","scale"), trControl=control.svmRadial)
> # summarize fit
> print(fit.svmRadial$results) #accuracy = 0.9231731
  sigma C Accuracy Kappa AccuracySD KappaSD
1 0.04419447 0.25 0.8672520 0.2457674 0.01803261 0.1117821
2 0.04419447 0.50 0.8996927 0.4850752 0.03357527 0.1936929
3 0.04419447 1.00 0.9026778 0.5334358 0.03377470 0.1934882
>
> #3. random forest
> control.rfe = trainControl(method="cv", number=5)
> set.seed(7)
> metric = "Accuracy"
> mtrv = 7 # mtrv=7 (number of variables to try)
```

2. Support Vector Machine

```
> set.seed(7)
> control.svmRadial = trainControl(method="cv", number=5)
> fit.svmRadial <- train(glyhb~, data=trainingData, method="svmRadial",
metric="Accuracy", preProc=c("center","scale"), trControl=control.svmRadial)
> # summarize fit
> print(fit.svmRadial$results)
```

```

RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins Project: (None)

Console Terminal Jobs
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/

> # split the data into training and validation sets
> set.seed(7)
> validation_index = createDataPartition(dataset$glyhb, p=0.90, list=FALSE)
> validationData = dataset[-validation_index,]
> trainingData = dataset[validation_index,]
>
> # comparison among different classifiers
> #1. Logistic Regression
> set.seed(7)
> control.glm = trainControl(method = "cv", number = 5)
> fit.glm = train(glyhb~, data = trainingData, method = "glm", preProc = c("center", "scale"), trControl = control.glm)
> print(fit.glm$results) # accuracy = 0.9172205
parameter Accuracy Kappa AccuracySD KappaSD
1 none 0.9113696 0.6246138 0.03189889 0.1209417
>
> #2. Support Vector Machine
> set.seed(7)
> control.svmRadial = trainControl(method="cv", number=5)
> fit.svmRadial <- train(glyhb~, data=trainingData, method="svmRadial", metric="Accuracy", preProc=c("center", "scale"), trControl=control.svmRadial)
> # summarize fit
> print(fit.svmRadial$results) #accuracy = 0.9231731
sigma C Accuracy Kappa AccuracySD KappaSD
1 0.04419447 0.25 0.8672520 0.2457674 0.01803261 0.1117821
2 0.04419447 0.50 0.8996927 0.4850752 0.03357527 0.1936929
3 0.04419447 1.00 0.9026778 0.5334358 0.03377470 0.1934882
>
> #3. random forest
> control.rf = trainControl(method="cv", number=5)
> set.seed(7)
> metric = "Accuracy"
> mtry = 7 # mtry=7 (number of variables to try)

```

3. Random forest

```

> control.rf = trainControl(method="cv", number=5)
> set.seed(7)
> metric = "Accuracy"
> mtry = 7 # mtry=7 (number of variables to try)
> tuneGrid <- expand.grid(.mtry=mtry)
> fit.rf_default <- train(glyhb~, data=trainingData, method="rf", metric=metric,
tuneGrid=tuneGrid, preProc=c("center", "scale"), trControl=control)
> print(fit.rf_default$results)

```

```

RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins Project: (None)

Console Terminal Jobs
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/

1 0.04419447 0.25 0.8672520 0.2457674 0.01803261 0.1117821
2 0.04419447 0.50 0.8996927 0.4850752 0.03357527 0.1936929
3 0.04419447 1.00 0.9026778 0.5334358 0.03377470 0.1934882
>
> #3. random forest
> control.rf = trainControl(method="cv", number=5)
> set.seed(7)
> metric = "Accuracy"
> mtry = 7 # mtry=7 (number of variables to try)
> tuneGrid <- expand.grid(.mtry=mtry)
> fit.rafo_default <- train(glyhb~, data=trainingData, method="rf", metric=metric, tuneGrid=tuneGrid, preProc=c("center", "scale"), trControl=control.rf)
> print(fit.rafo_default$results) # accuracy = 0.9114924
mtry Accuracy Kappa AccuracySD KappaSD
1 7 0.9262072 0.6997773 0.02769387 0.1060189
>
> #4. parameter tuning via grid search for random forest
> control.rf_search <- trainControl(method="repeatedcv", number=6, repeats=3, search="grid")
> set.seed(7)
> tuneGrid <- expand.grid(.mtry=c(1:18))
> fit.rf_gridsearch <- train(glyhb~, data=trainingData, method="rf", metric=metric, tuneGrid=tuneGrid, trControl=control.rf_search, ntree=1000)
There were 50 or more warnings (use warnings() to see the first 50)
> print(fit.rf_gridsearch) # accuracy = 0.9204355 when mtry = 12
Random Forest

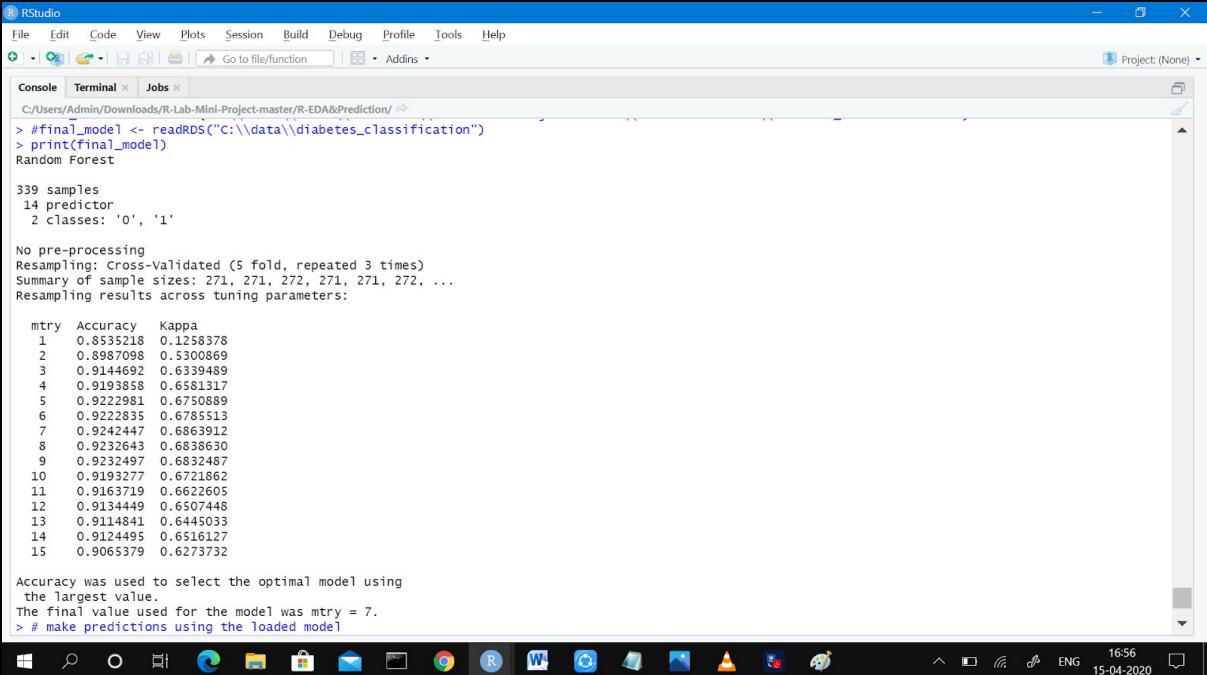
339 samples
14 predictor
2 classes: '0', '1'

No pre-processing
Resampling: Cross-Validated (6 fold, repeated 3 times)
Summary of sample sizes: 282, 282, 282, 283, 283, 283, ...

```

> #4. Parameter tuning via grid search for random forest

```
> control.rf_search <- trainControl(method="repeatedcv", number=5,
repeats=3, search="grid")
> set.seed(7)
> tuneGrid <- expand.grid(.mtry=c(1:15))
> fit.rf_gridsearch <- train(glyhb~., data=trainingData, method="rf",
metric=metric, tuneGrid=tuneGrid, trControl=control.rf_search, ntree=1000)
> print(fit.rf_gridsearch) # accuracy = 0.9204355 when mtry = 12
```



```
> #final_model <- readRDS("C:\\data\\diabetes_classification")
> print(final_model)
Random Forest

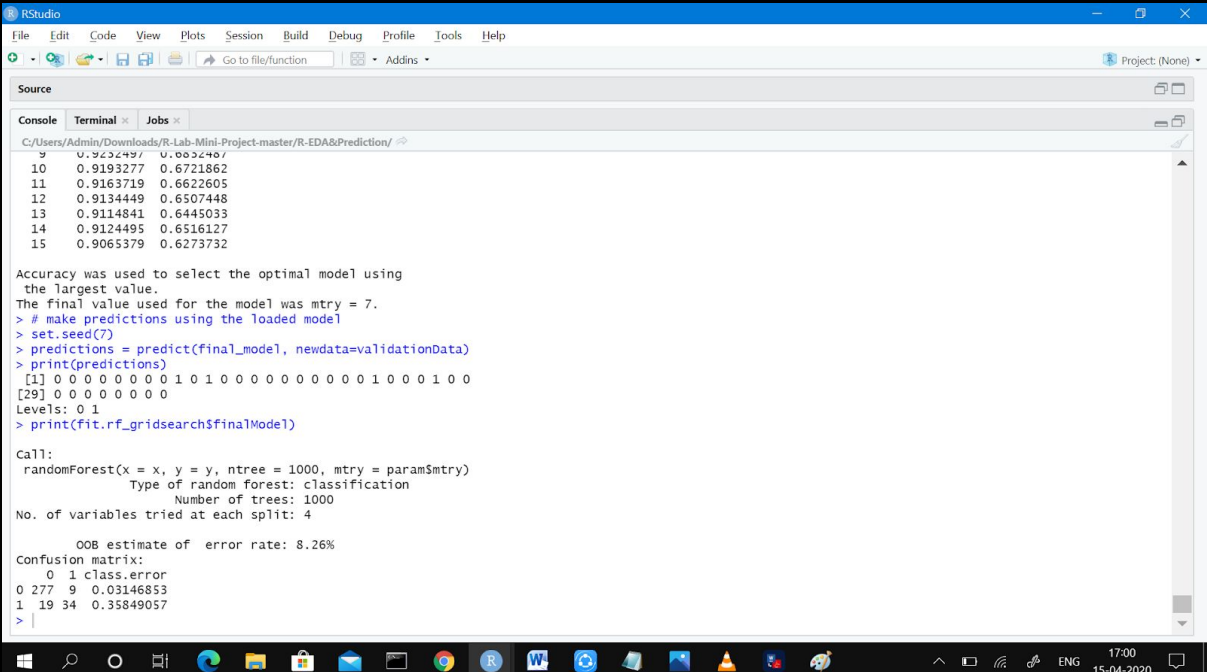
339 samples
14 predictor
2 classes: '0', '1'

No pre-processing
Resampling: Cross-Validated (5 fold, repeated 3 times)
Summary of sample sizes: 271, 271, 272, 271, 272, ...
Resampling results across tuning parameters:

mtry Accuracy Kappa
1 0.8535218 0.1258378
2 0.8987098 0.5300869
3 0.9144692 0.6339489
4 0.9193858 0.6581317
5 0.9222981 0.6750889
6 0.9222835 0.6785513
7 0.9242447 0.6863912
8 0.9232643 0.6838630
9 0.9232497 0.6832487
10 0.9193277 0.6721862
11 0.9163719 0.6622605
12 0.9134449 0.6507448
13 0.9114841 0.6445033
14 0.9124495 0.6516127
15 0.9065379 0.6273732

Accuracy was used to select the optimal model using
the largest value.
The final value used for the model was mtry = 7.
> # make predictions using the loaded model
```

```
> print(fit.rf_gridsearch$finalModel)
```



```
10 0.9193277 0.6721862
11 0.9163719 0.6622605
12 0.9134449 0.6507448
13 0.9114841 0.6445033
14 0.9124495 0.6516127
15 0.9065379 0.6273732


Accuracy was used to select the optimal model using
the largest value.
The final value used for the model was mtry = 7.
> # make predictions using the loaded model
> set.seed(7)
> predictions = predict(final_model, newdata=validationData)
> print(predictions)
[1] 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0
[29] 0 0 0 0 0 0 0 0
Levels: 0 1
> print(fit.rf_gridsearch$finalModel)

Call:
randomForest(x = x, y = y, ntree = 1000, mtry = param$mtry)
Type of random forest: classification
Number of trees: 1000
No. of variables tried at each split: 4

OOB estimate of error rate: 8.26%
Confusion matrix:
 0 1 class.error
0 277 9 0.03146853
1 19 34 0.35849057
>
```

```
> plot(fit.rf_gridsearch)
> # make predictions on the validation set
> set.seed(7)
> predictions = predict(fit.rf_gridsearch, newdata=validationData)
> confusionMatrix = confusionMatrix(predictions, validationData$glyhb)
> # confusion matrix
> print(confusionMatrix$table)
```

```
> # make predictions on the validation set
> set.seed(7)
> predictions = predict(fit.rf_gridsearch, newdata=validationData)
> confusionMatrix = confusionMatrix(predictions, validationData$glyhb)
> # confusion matrix
> print(confusionMatrix$table)
      Reference
Prediction 0 1
      0 31 3
      1  0 2
```



PART 4: Prediction of Accuracy

#save the final classifier model into disk

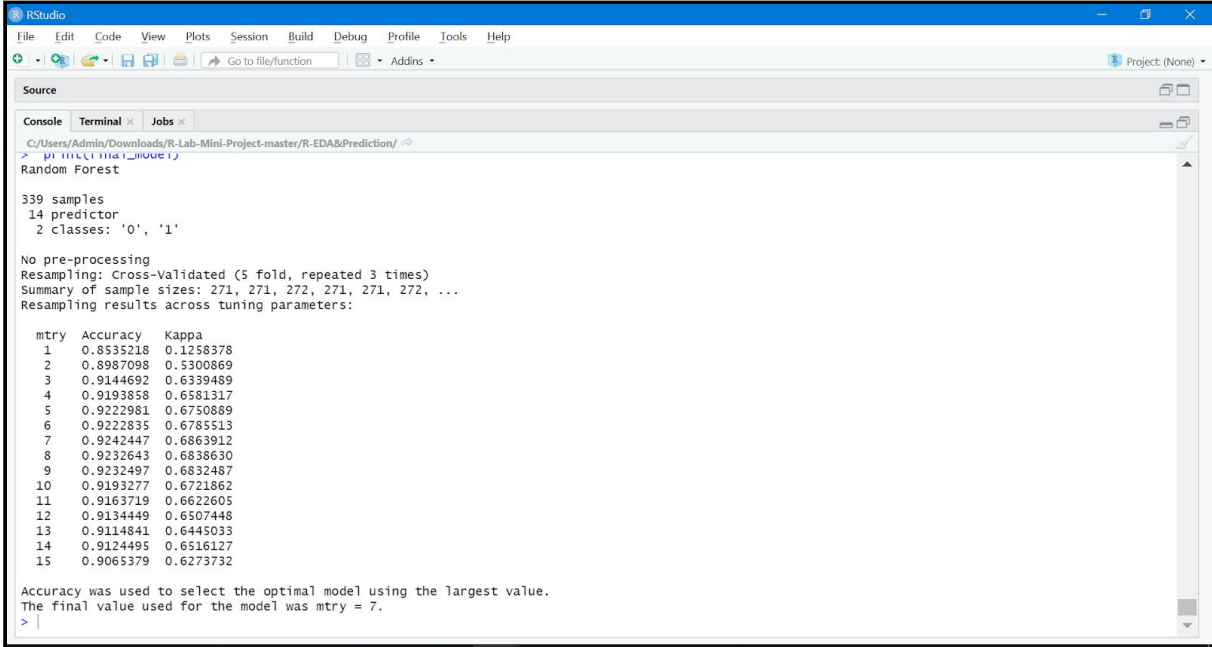
```
> saveRDS(fit.rf_gridsearch, "C:\\Users\\Admin\\Desktop\\  
R-Lab-Mini-Project\\diabetes_classification1")
```

> # load the model from the disk

```
> final_model <- readRDS("C:\\Users\\Admin\\Desktop\\R-Lab-Mini-Project  
\\diabetes_classification")
```

```
> #final_model <- readRDS("C:\\data\\diabetes_classification")
```

```
> print(final_model)
```



```
C:/Users/Admin/Downloads/R-Lab-Mini-Project-master/R-EDA&Prediction/ > print(fit.rf_gridsearch)
Random Forest

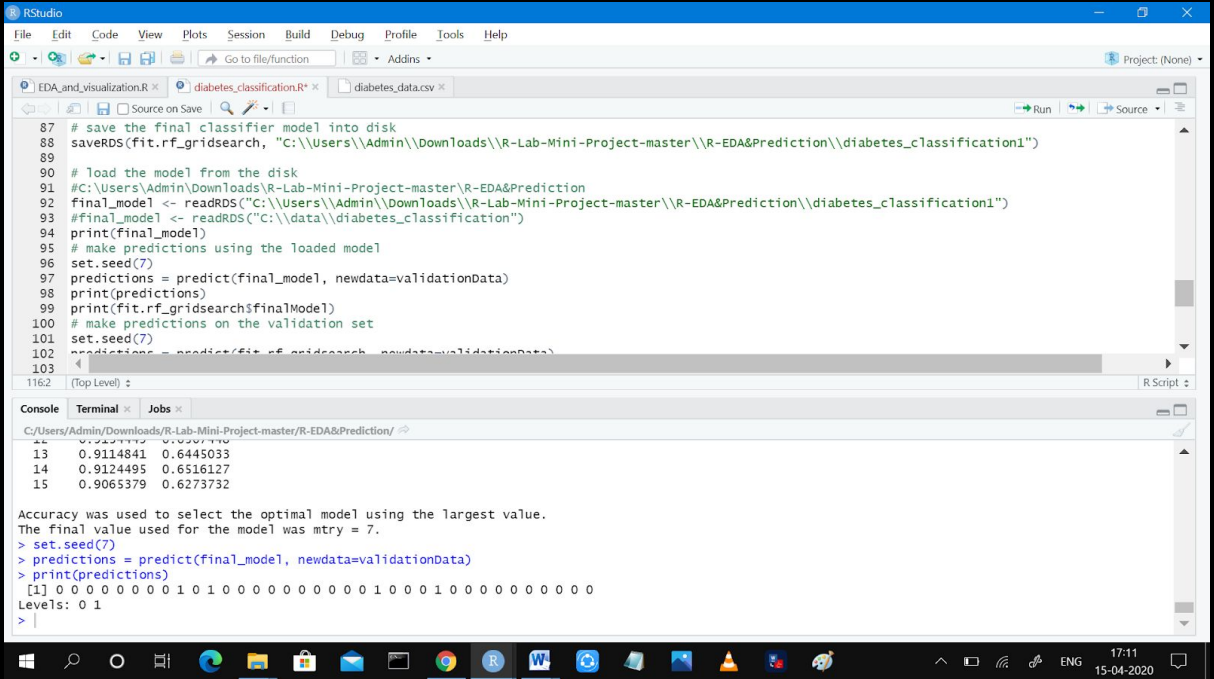
339 samples
14 predictor
2 classes: '0', '1'

No pre-processing
Resampling: Cross-Validated (5 fold, repeated 3 times)
Summary of sample sizes: 271, 271, 272, 271, 271, 272, ...
Resampling results across tuning parameters:

mtry  Accuracy  Kappa
1    0.8535218  0.1258378
2    0.8987098  0.5300869
3    0.9144692  0.6339489
4    0.9193858  0.6581317
5    0.9222981  0.6750889
6    0.9222835  0.6785513
7    0.9242447  0.6863912
8    0.9232643  0.6838630
9    0.9232497  0.6832487
10   0.9193277  0.6721862
11   0.9163719  0.6622605
12   0.9134449  0.6507448
13   0.9114841  0.6445033
14   0.9124495  0.6516127
15   0.9065379  0.6273732

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was mtry = 7.
> |
```

```
# make predictions using the loaded model
> set.seed(7)
> predictions = predict(final_model, newdata=validationData)
> print(predictions)
```



The screenshot shows the RStudio interface with three open files: `EDA_and_visualization.R`, `diabetes_classification.R`, and `diabetes_data.csv`. The `diabetes_classification.R` file contains the following code:

```
87 # save the final classifier model into disk
88 saveRDS(fit.rf_gridsearch, "C:\\Users\\Admin\\Downloads\\R-Lab-Mini-Project-master\\R-EDA&Prediction\\diabetes_classification1")
89
90 # load the model from the disk
91 #C:\\Users\\Admin\\Downloads\\R-Lab-Mini-Project-master\\R-EDA&Prediction
92 final_model <- readRDS("C:\\Users\\Admin\\Downloads\\R-Lab-Mini-Project-master\\R-EDA&Prediction\\diabetes_classification1")
93 #final_model <- readRDS("C:\\data\\diabetes_classification")
94 print(final_model)
95 # make predictions using the loaded model
96 set.seed(7)
97 predictions = predict(final_model, newdata=validationData)
98 print(predictions)
99 print(fit.rf_gridsearch$finalModel)
100 # make predictions on the validation set
101 set.seed(7)
102 predictions = predict(fit.rf_gridsearch, newdata=validationData)
103
```

The console output shows the following results:

```
13 0.9114841 0.6445033
14 0.9124495 0.6516127
15 0.9065379 0.6273732
```

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was mtry = 7.

```
> set.seed(7)
> predictions = predict(final_model, newdata=validationData)
> print(predictions)
[1] 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0
Levels: 0 1
```